

1) From the Groq and Openairouter : List out the multi models

Groq:

Models :

1) LLaVA (Large Language and Vision Assistant)

Example: llava-v1.5-7b

Supports: Text + Image input

2) Gemma Vision (if available via Groq)

Supports: Text + Image

OpenRouter:

Models:

1) GPT-4o

o Supports: Text + Image + Audio

o Fully multimodal (most powerful)

2) GPT-4o mini

o Supports: Text + Image

o Faster and cheaper version

3) DALL·E 3

o Supports: Text → Image

4) GPT-4o Transcribe

.

Supports : Audio +Text

5) gpt-realtime-mini

A cost-efficient version of GPT Realtime

Supports : Text , Image and audio

6) GPT-5.4

Supports : Text + Image

2. Take 3 models from each service provider

Context window, Input output token cost

1) Chat gpt

Model : gpt 5.4

context window - 1,050,000

Input : Text , Image

Output : Text

Token cost : Input - \$2.5 , coached input - \$ 0.075

and Output - \$4.50 per 1M tokens

Model :GPT-5.4 mini

context window - 400,000

Input : Text , Image

Output : Text

Token cost : Input - \$0.75 , coached input - \$ 0.25

and Output - \$15.0 per 1M tokens

Model: GPT-4o Transcribe

context window - 16,000

Input : Text , Audio

Output : Text

Token cost : Input - \$2.50 , Output - \$10.00 per 1M

tokens

2) Gemini

Model: gemini-3.1-pro

Context Window: 1,048,576 (Up to 2,000,000 for specific

enterprise tiers)

Input: Text, Image, Video, Audio, PDF

Output: Text (66K token limit)

Token Cost: Input - \$2.00 , Cached Input - \$0.20 and
Output - \$12.00 per 1M tokens
Model: Gemini 3 Flash
Models: gemini-3-flash
Context Window: 1,048,576
Input: Text, Image, Video, Audio
Output: Text (66K token limit)
Token Cost: Input - \$0.50 , Cached Input - \$0.05 and
Output - \$3.00 per 1M tokens
Model: Gemini 2.5 Flash
Models: gemini-2.5-flash
Context Window: 1,048,576
Input: Text, Image, Video
Output: Text (8K token limit)
Token Cost: Input - \$0.30 , Cached Input - \$0.03 and
Output - \$2.50 per 1M tokens
3) Deepseek
Model: DeepSeek-V3.2
Context Window: 128,000 tokens
Input: Text
Output: Text (4K default, 8K maximum)
Token Cost: Input - \$0.28 , Cached Input - \$0.028 and
Output - \$0.42 per 1M tokens
Model: DeepSeek-V3.2 (Thinking Mode)
Context Window: 128,000 tokens
Input: Text
Output: Text (32K default, 64K maximum)
Token Cost: Input - \$0.28 , Cached Input - \$0.028 and
Output - \$0.42 per 1M tokens
Model: DeepSeek-R1
Context Window: 128,000 tokens
Input: Text
Output: Text (16,384 tokens maximum)
Token Cost: Input - \$0.14 (cache hit), \$0.55 (cache
miss) and Output - \$2.19 per 1M
tokens

3) Use the Tokenizer with different prompt and see how the tokens are generated

1) write test case for Spotify Login functionality

Tokens - 8

Characters - 48

Token id's : [9566, 1746, 1890, 395, 53142, 14900, 22622, 220]

2) Give me some scenarios for banking applications

Tokens - 7

Characters - 47

Token id's : [50117, 668, 1236, 42318, 395, 31678, 9391]

3) how to identify device discovery in Spotify mobile app

Tokens - 9

Characters - 54

Token id's : [8923, 316, 11410, 5897, 30465, 306, 53142, 7141, 1053]